

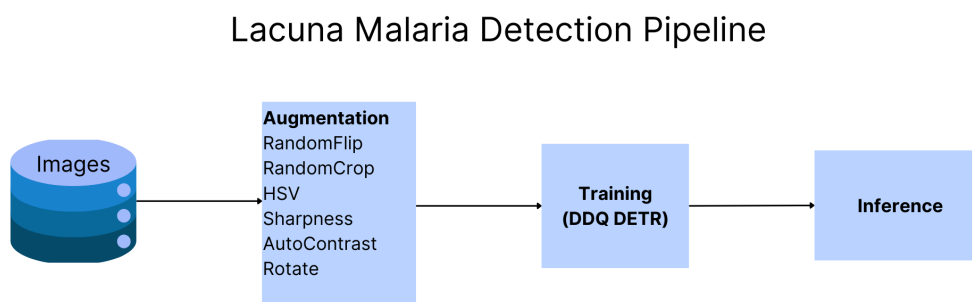
# Lacuna Malaria Detection Challenge - 3rd Place solution

## 1. Overview

The objective of the challenge was to develop a multiclass object detection and classification model capable of accurately localizing and classifying malaria parasites in blood slide images.

The solution employs a single DDQ DETR model, trained on the complete training dataset using the MMDetection library.

## 2. Architecture diagram



## 3. ETL process

The data was extracted to mirror the directory structure below:

```
.
├── /workspace/
│   ├── mmdetection/
│   │   ├── data/
│   │   │   ├── images/
│   │   │   ├── Test.csv
│   │   │   ├── Train.csv
│   │   │   └── SampleSubmission.csv
```

The data was transformed into the COCO annotation format required by the MMDetection library. This included mapping bounding boxes and class labels to COCO-compliant JSON files.

## 4. Data modeling

The following augmentations were applied to the training dataset during training:

- *RandomFlip*
- *RandomChoiceResize*
- *YOLOXHSVRandomAug*
- *Sharpness*
- *AutoContrast*
- *Rotation*

The Swin-L variant of the Dense Distinct Query for End-to-End Object Detection (DDQ-DETR) model was selected based on its promising performance on object detection tasks. The model was trained for 30 epochs with the following parameters:

### Training Parameters

- **Epochs:** 30
- **Batch size:** 2
- **Optimizer:** AdamW (learning rate = 0.0002, weight decay = 0.05)
- **LR Schedule:**
  - **Base LR:** 1e-4
  - **Warmup:** LinearLR warm-up for the first 2000 iterations, increasing learning rate from  $0.0001 \times \text{base\_lr}$  to  $\text{base\_lr}$ .
  - **Multi-step Decay:** Constant learning rate until epoch 20, reduced by  $\gamma = 0.1$  at epochs 20 and 26.

## 5. Inference

The model was deployed using the MMDetection framework on a GPU-enabled machine for efficient inference.

Deployment workflow:

1. Load trained model weights (`epoch_30.pth`) and configuration file (`cfg`) with `init_detector`.
2. Sequentially process test images from a DataFrame (`df_test`) containing image paths and IDs.
3. Perform object detection using `inference_detector` to generate bounding boxes, class labels, and confidence scores.

#### 4. Apply post-processing:

- Filter predictions with a confidence threshold (`minconf = 0.05`).
- Record a default entry with a class label of `NEG` if no detections are found.
- Structure and save detected bounding boxes, scores, and labels in a dictionary format.

#### Model updates and retraining:

- New data can be incorporated through fine-tuning or full retraining.
- If fine-tuning, modify the `load_from` parameter in the configuration to use the current model weights.

### 6. Run time

- Training: 12 hours
- Inference: 15 minutes

### 7. Performance metrics

- **Disk space:** The model required ~100 GB to store checkpoints after each epoch during training.
- **GPU VRAM usage:** Peak usage was ~45 GB during training.
- **Evaluation scores:**
  - **Public leaderboard mAP@0.5:** 0.92529841
  - **Private leaderboard mAP@0.5:** 0.921717525